

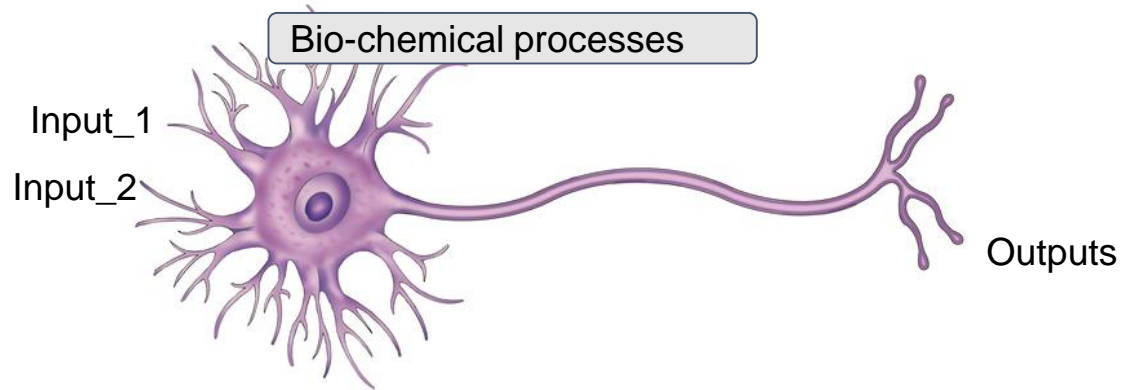
# Deep Learning Fundamentals

(A crash course)

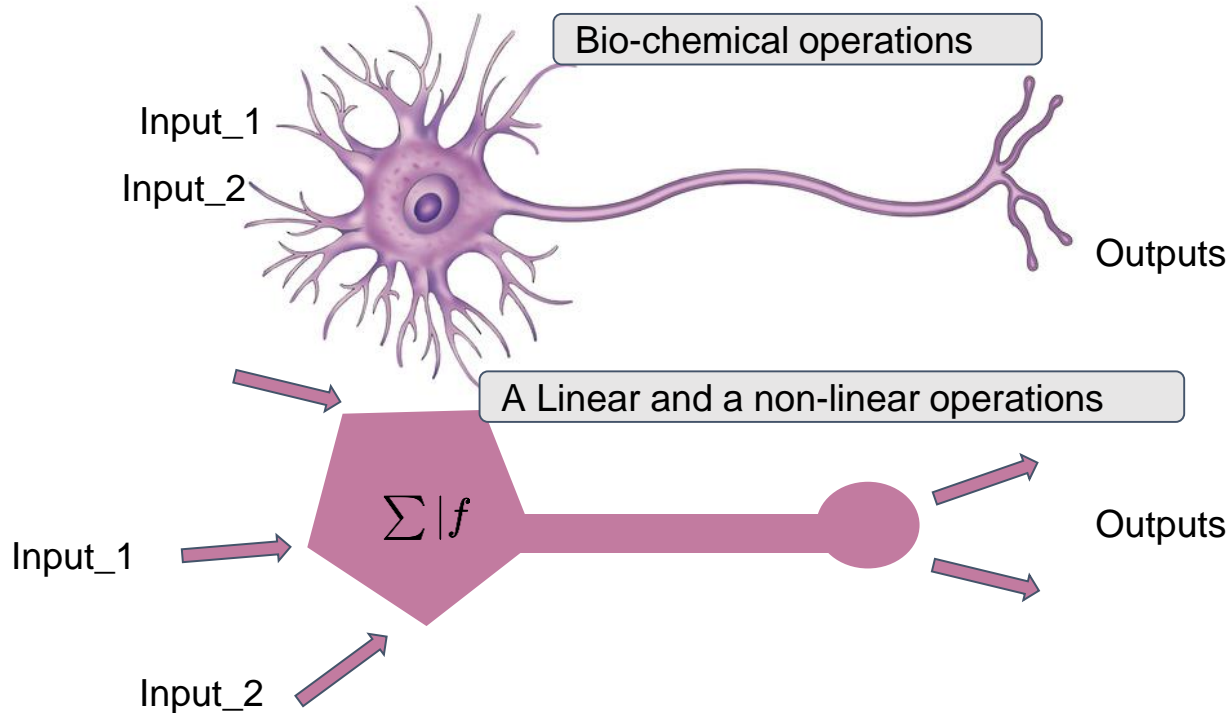
Danda Pani Paudel  
**Computer Vision Lab @ ETH Zurich**



# Biological Neurons and Artificial Neurons

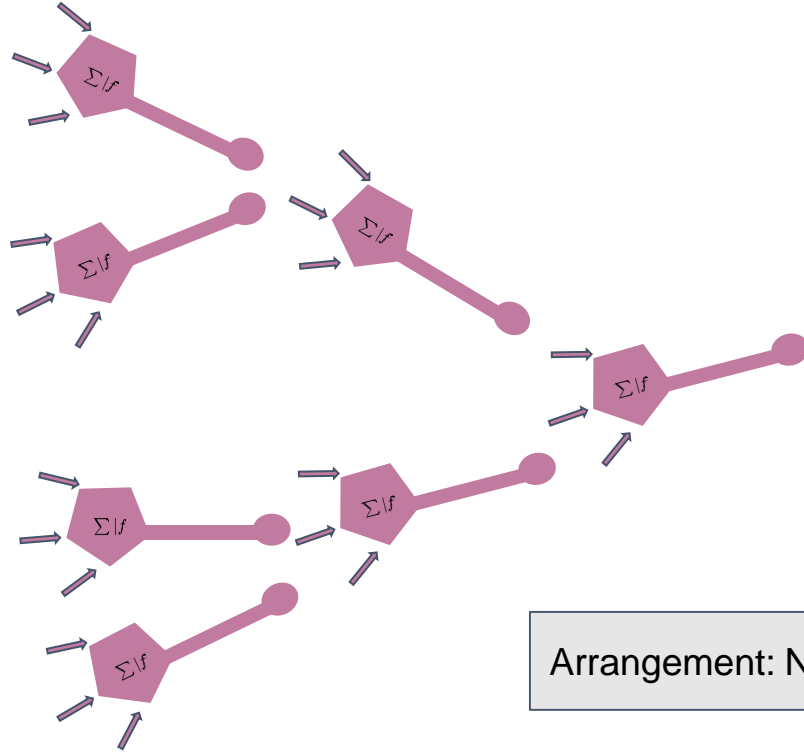


# Biological Neurons and Artificial Neurons



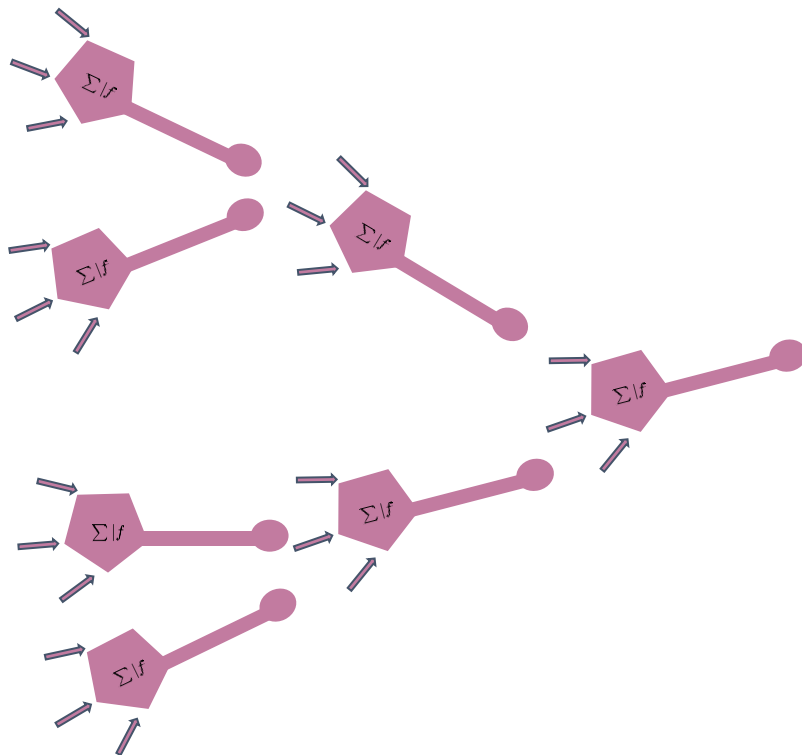


# Many Artificial Neurons



Arrangement: Network Architecture

# What is so special about it?



Parameters

Data/Examples

Input data

Output labels



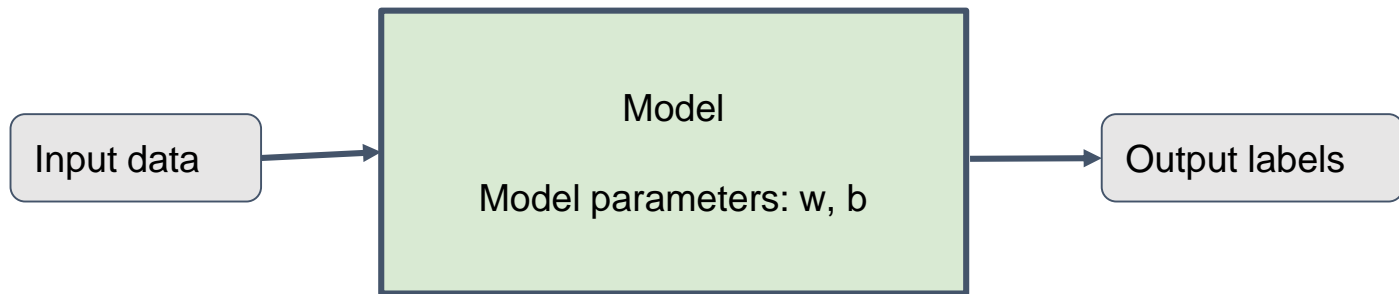






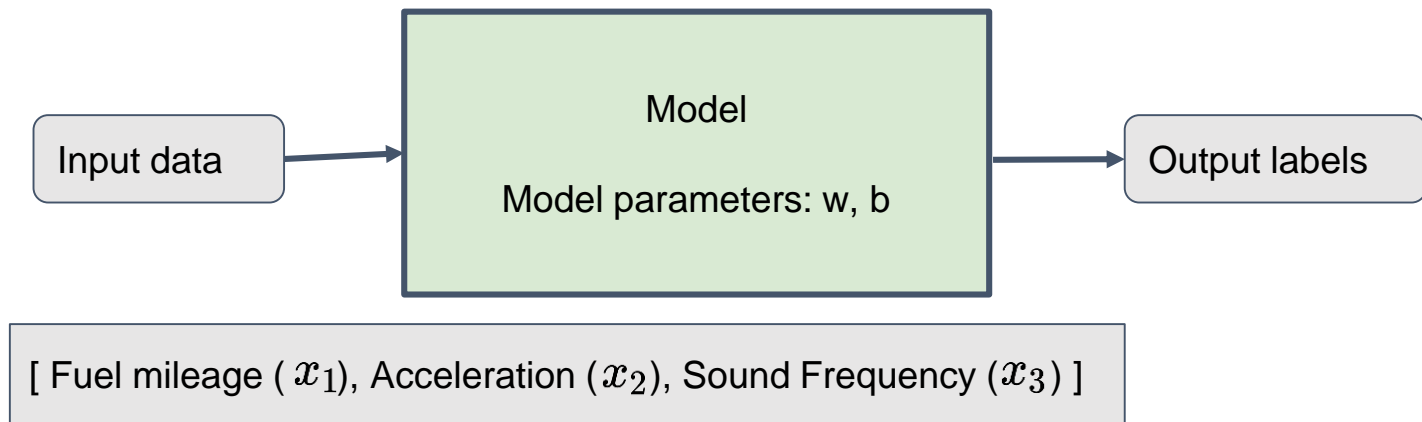
# Logistic regression

Binary classification



# Logistic regression

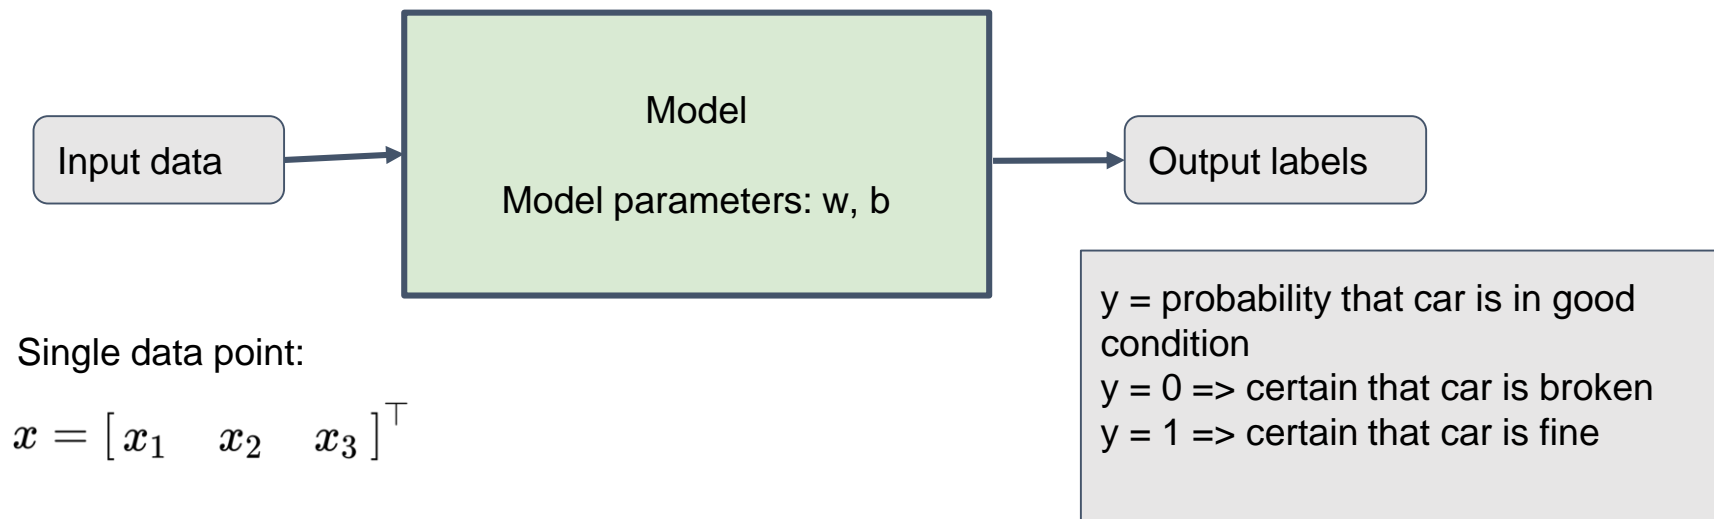
## Binary classification example



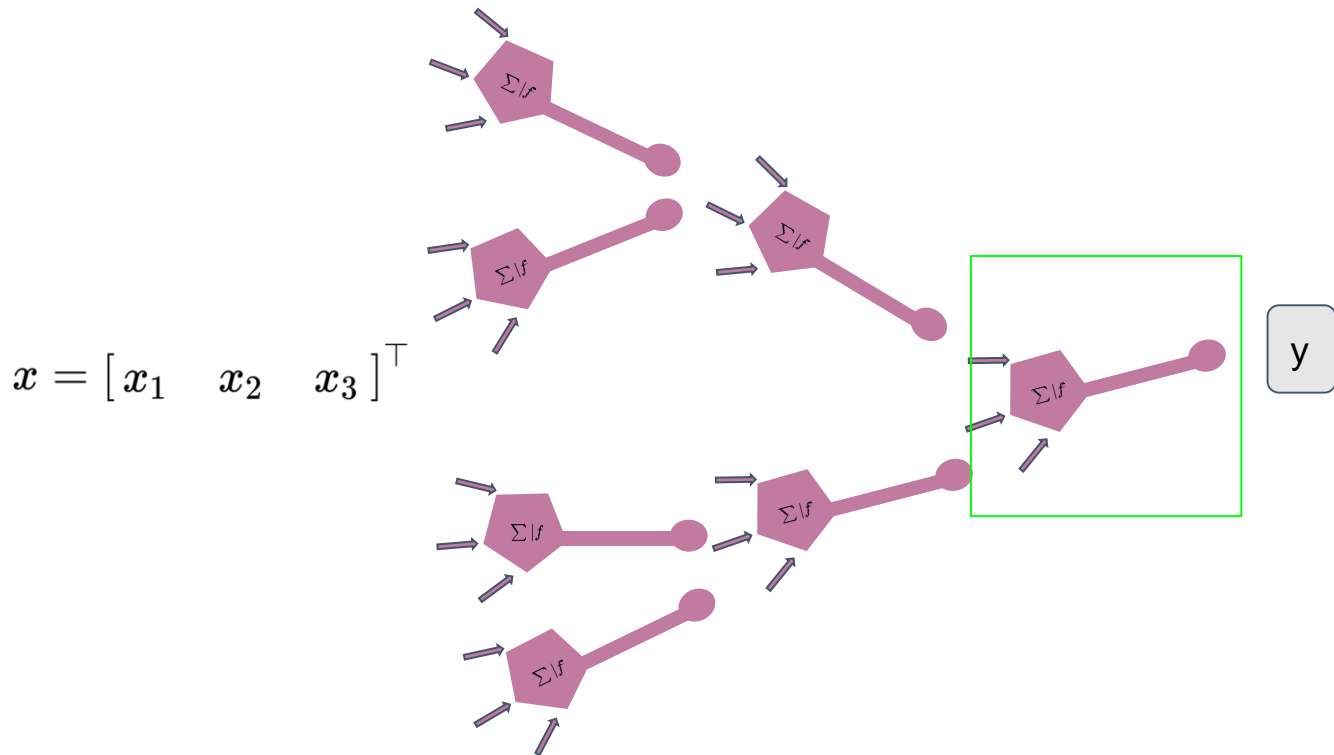
$$\mathbf{x} = [x_1 \quad x_2 \quad x_3]^\top$$

# Logistic regression

## Binary classification example

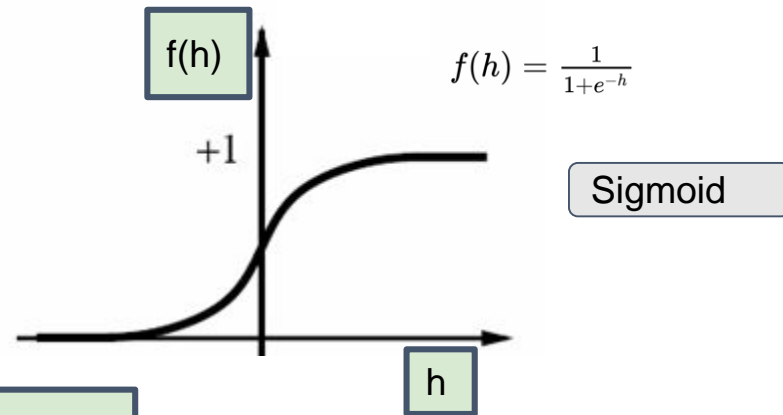
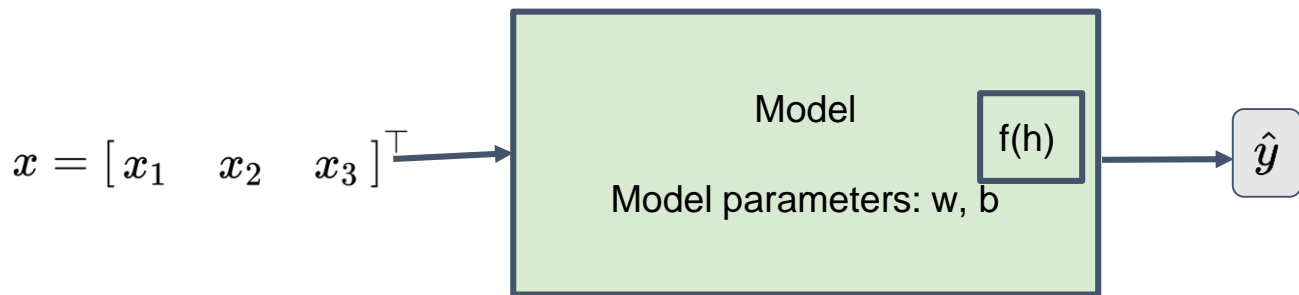


# Logistic Regression



# Logistic regression

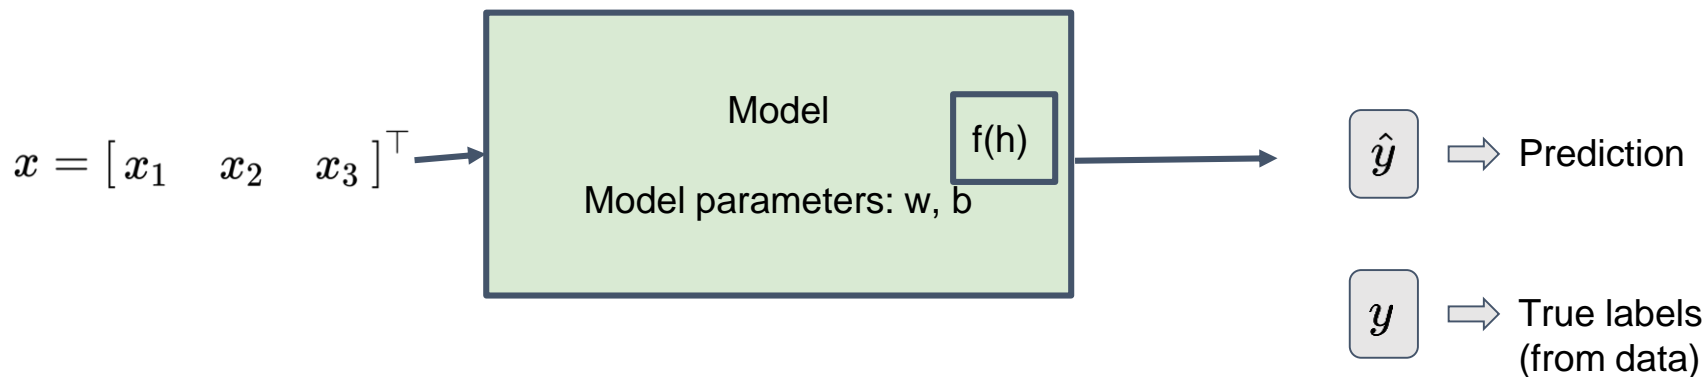
Binary classification example



Sigmoid squeezes the results between 0 and 1

# Logistic regression

## Binary classification example

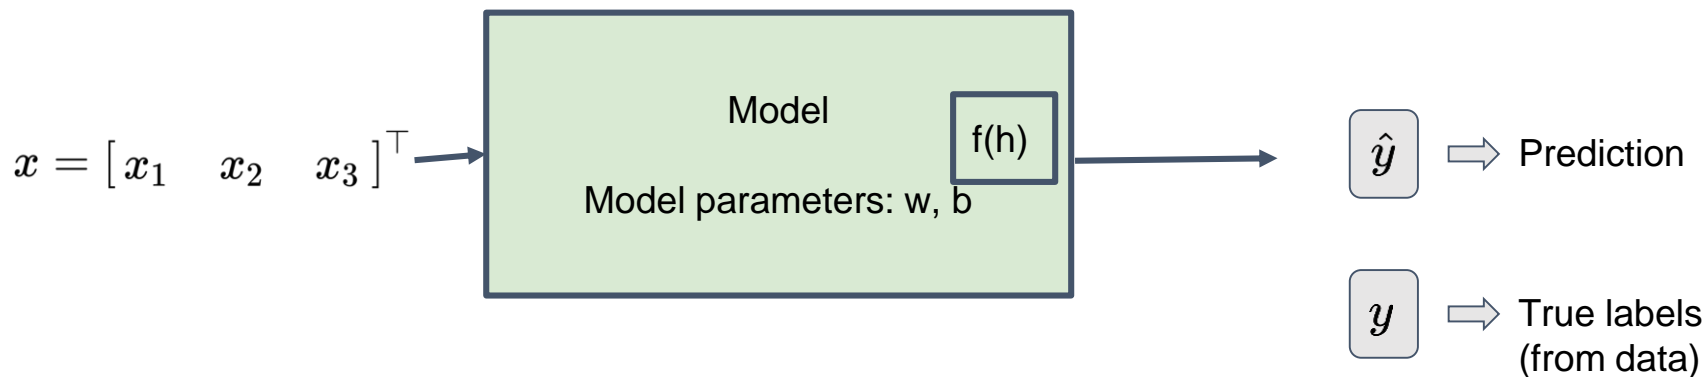


Logistic regression loss: Error of prediction from the true label

$$-(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$$

# Logistic regression

## Binary classification example



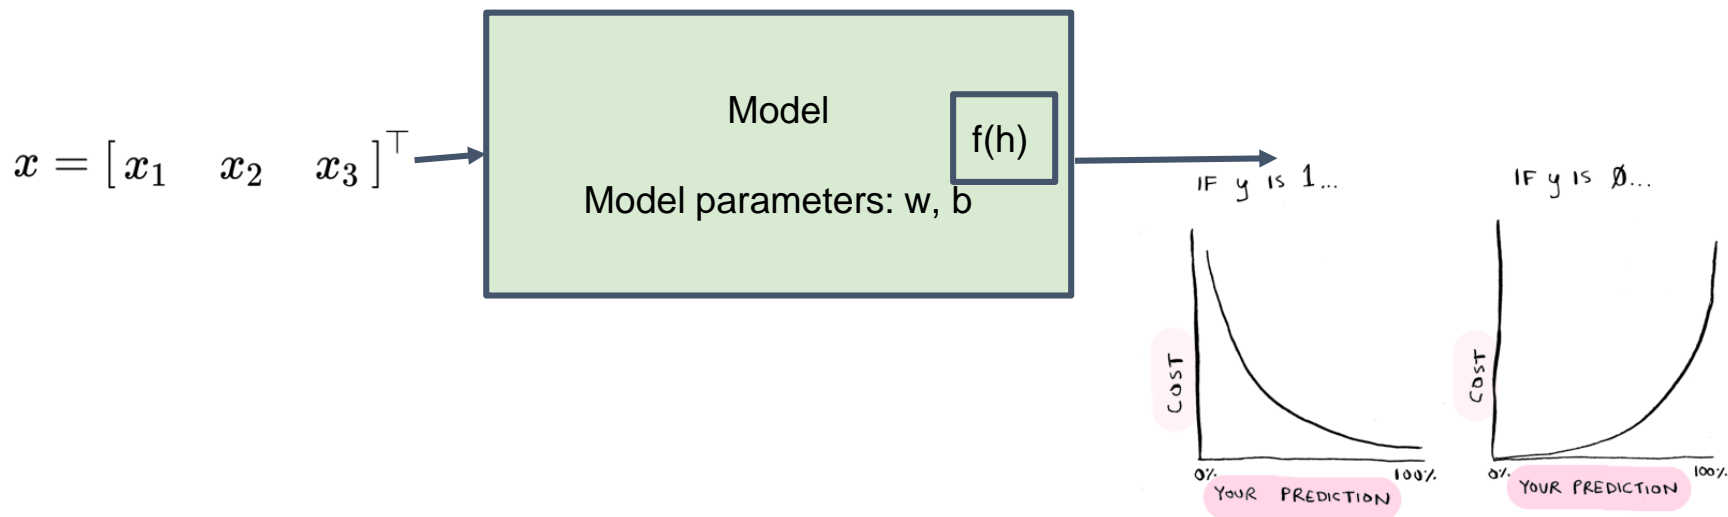
Logistic regression loss: Error of prediction from the true label

$$-(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$$



# Logistic regression

## Binary classification example



$$-(y \log \hat{y} + (1 - y) \log(1 - \hat{y})) \quad 0 \rightarrow \infty$$

**A Convex Function w.r.t. h!!!**

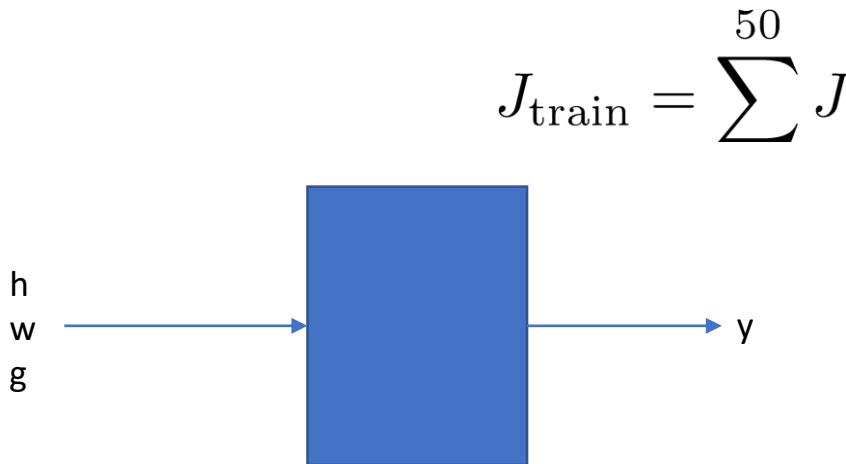
# Gradient Descent for Logistic Regression

## ❑ Dummy example

Training dataset:

50 people: each person is represented by **height, weight and gender**

50 people: predict ***overweight or not?***



## ❑ Gradient Descent evaluates all samples for each update

# Stochastic Gradient Descent

## ❑ Dummy example

$$J_{\text{sgd}} = J_i$$

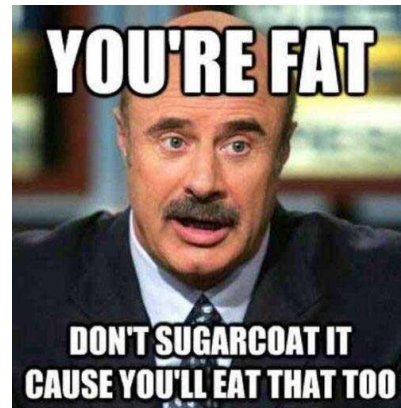
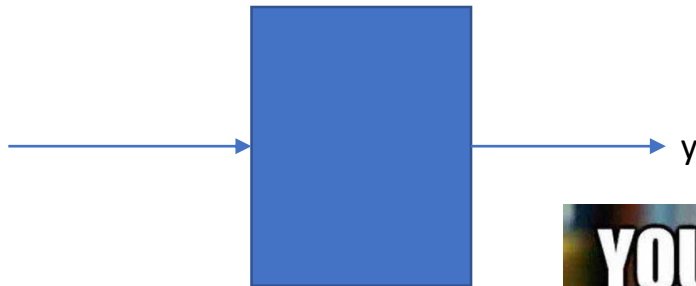
$i = \text{random\_index}$

Training dataset:

50 people: each person is represented by height, weight and gender

50 people: predict **overweight or not?**

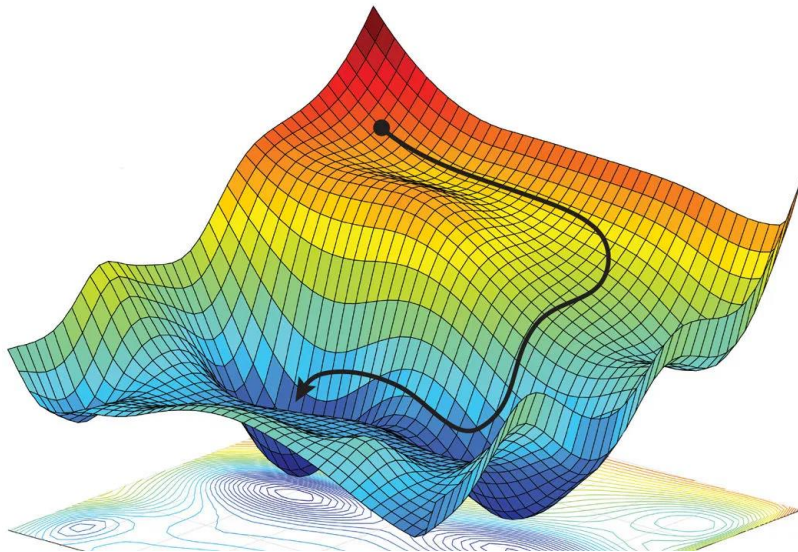
h  
w  
g



## ❑ SGD takes only one (few) random sample(s)

# Why Stochastic Gradient Descent?

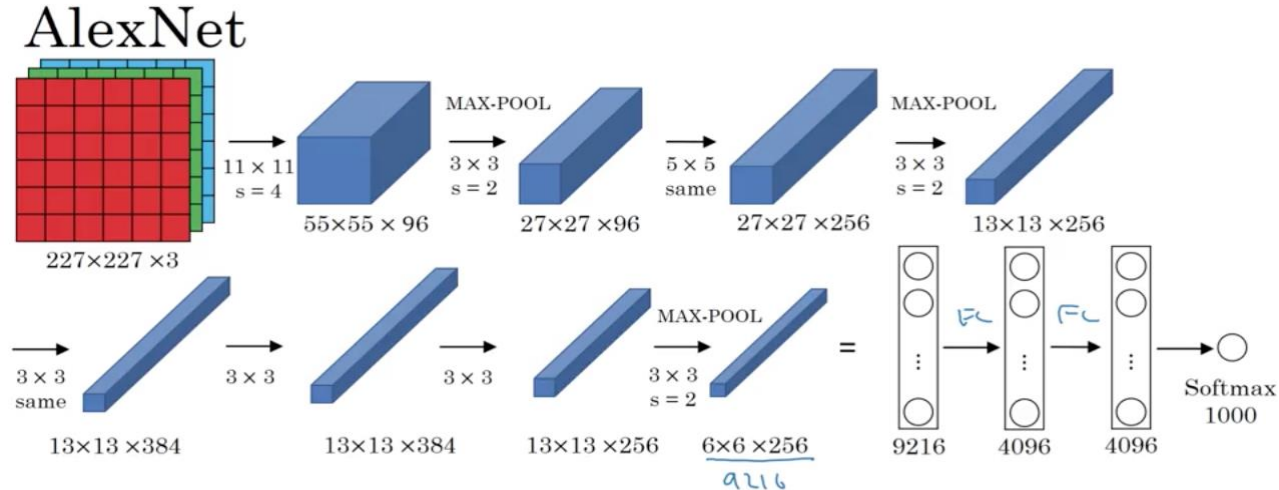
- Because of the memory/computational limitation
- To avoid local minima traps.... hopefully!



Source: <https://reconsider.news/2018/05/09/ai-researchers-allege-machine-learning-alchemy/>

# Sneak - peak into current deep learning

- ❑ ImageNet: Learn to classify images using **1M training examples**
- ❑ Using a deep neural Network of **60 M parameters!!**



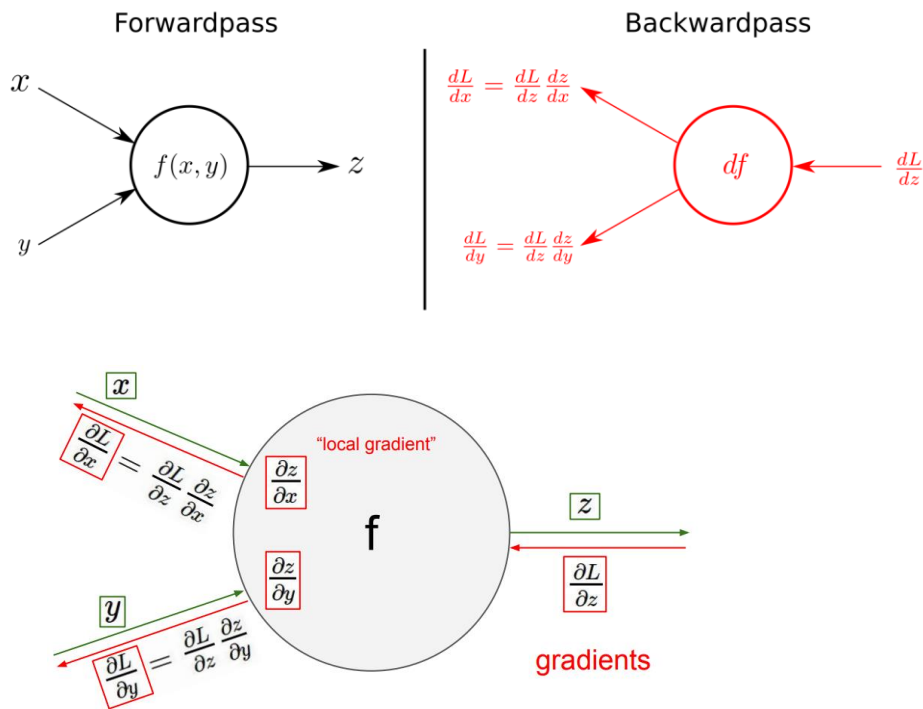
How do I optimize millions of parameters?



# Backpropagation

## Learning representations by back-propagating errors

David E. Rumelhart\*, Geoffrey E. Hinton† ©1986 Nature Publishing Group  
& Ronald J. Williams\*



$$E = \frac{1}{2} \sum_c \sum_j (y_{j,c} - d_{j,c})^2 \quad (3)$$

The backward pass starts by computing  $\partial E/\partial y$  for each of the output units. Differentiating equation (3) for a particular case,  $c$ , and suppressing the index  $c$  gives

$$\frac{\partial E}{\partial y_j} = y_j - d_j \quad (4)$$

We can then apply the chain rule to compute  $\partial E/\partial x_j$

$$\frac{\partial E}{\partial x_j} = \frac{\partial E}{\partial y_j} \cdot \frac{dy_j}{dx_j}$$

Differentiating equation (2) to get the value of  $dy_j/dx_j$  and substituting gives

$$\frac{\partial E}{\partial x_j} = \frac{\partial E}{\partial y_j} \cdot y_j(1 - y_j) \quad (5)$$

This means that we know how a change in the total input  $x$  to an output unit will affect the error. But this total input is just a linear function of the states of the lower level units and it is also a linear function of the weights on the connections, so it is easy to compute how the error will be affected by changing these states and weights. For a weight  $w_{ji}$ , from  $i$  to  $j$  the derivative is

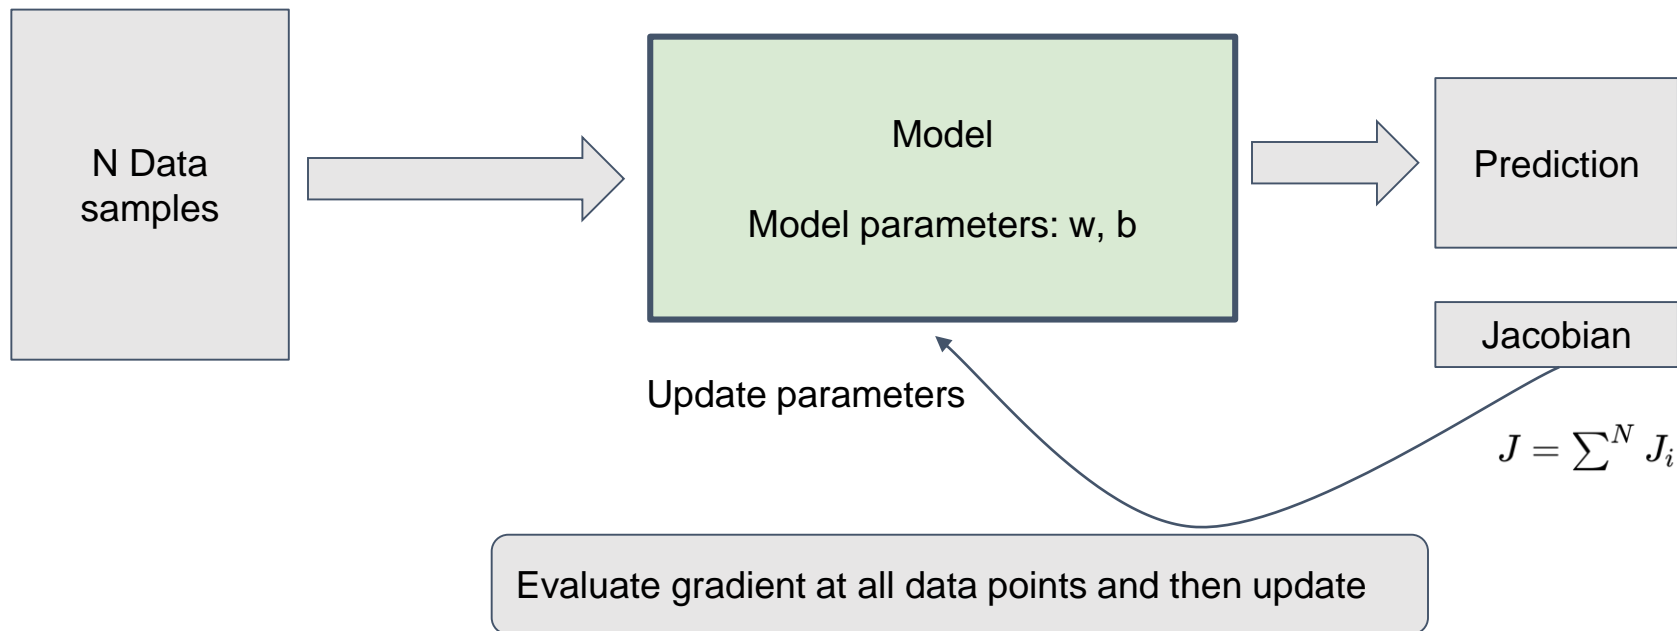
$$\begin{aligned} \frac{\partial E}{\partial w_{ji}} &= \frac{\partial E}{\partial x_j} \cdot \frac{\partial x_j}{\partial w_{ji}} \\ &= \frac{\partial E}{\partial x_j} \cdot y_i \end{aligned} \quad (6)$$





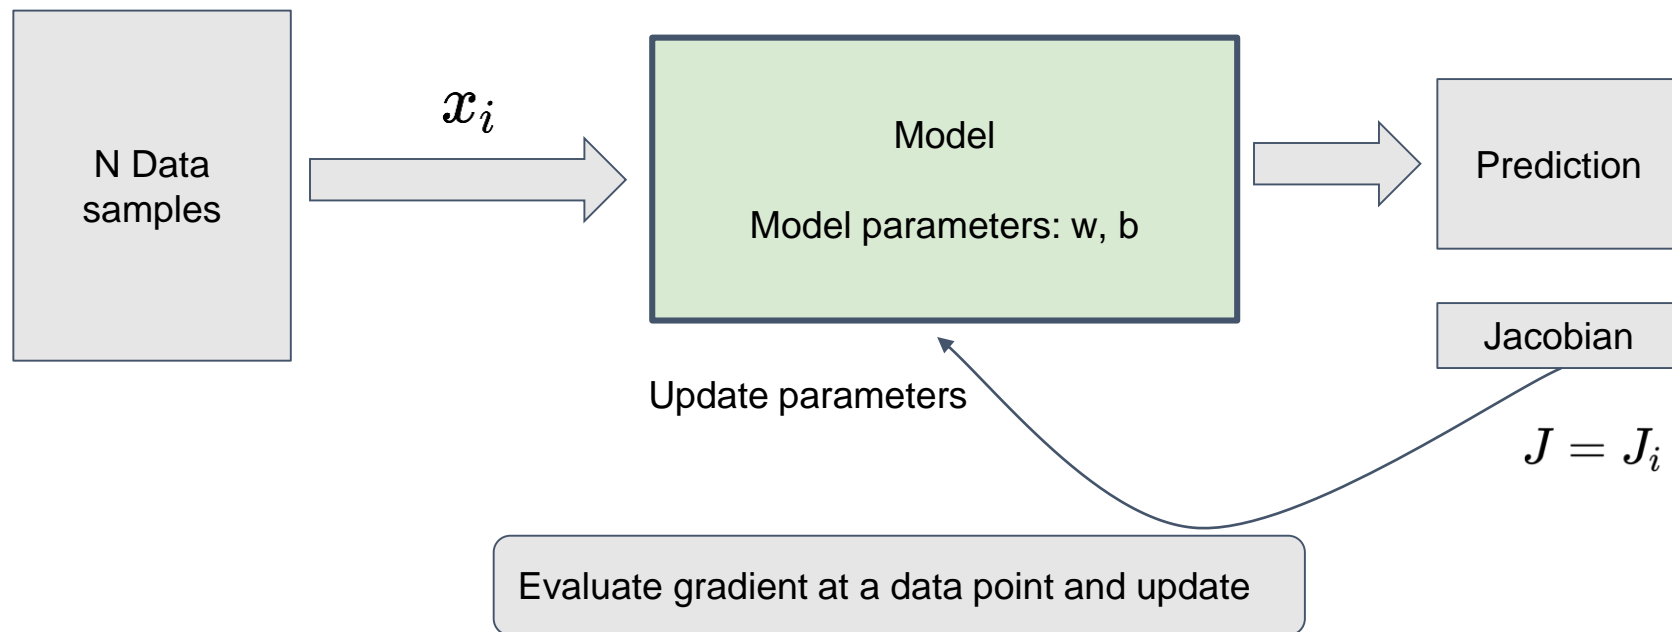
# Training in practice

## Gradient descent



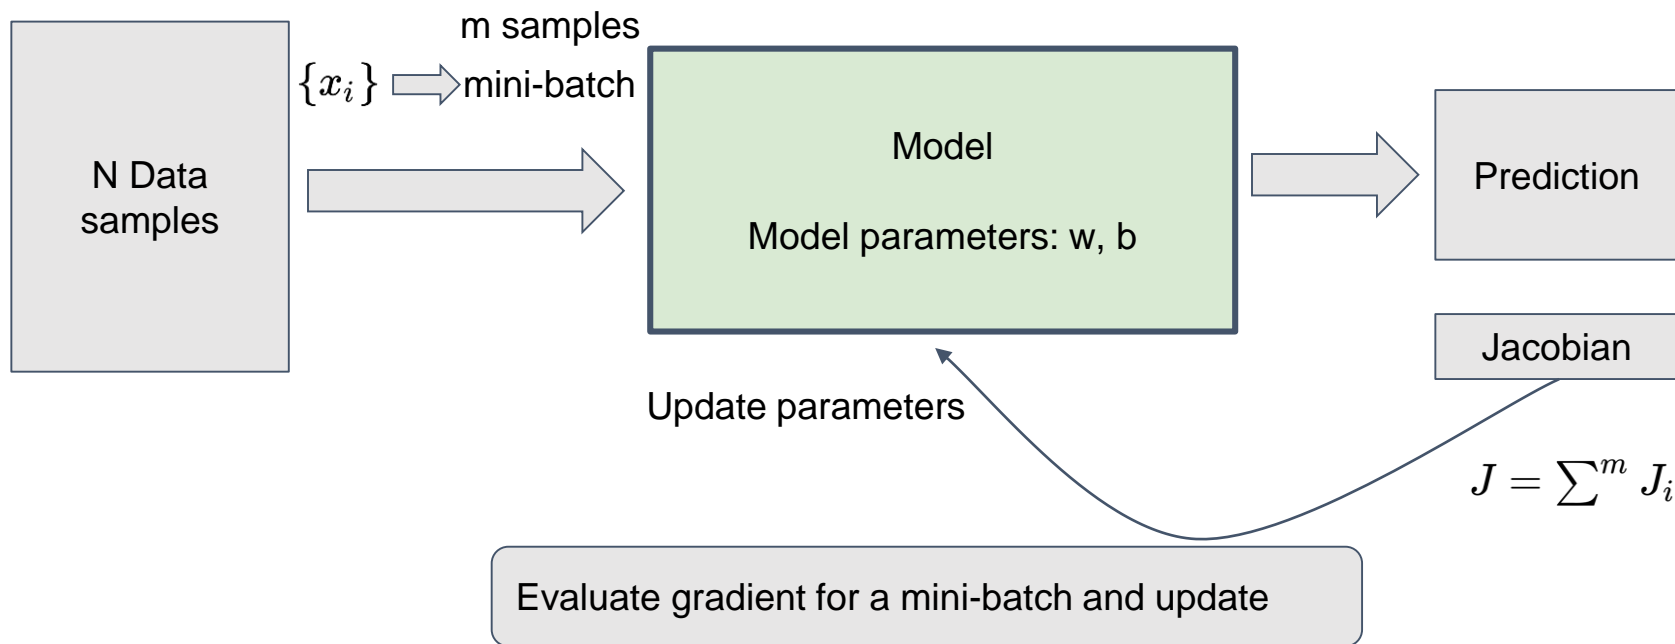
# Training in practice

## Stochastic Gradient descent



# Training in practice

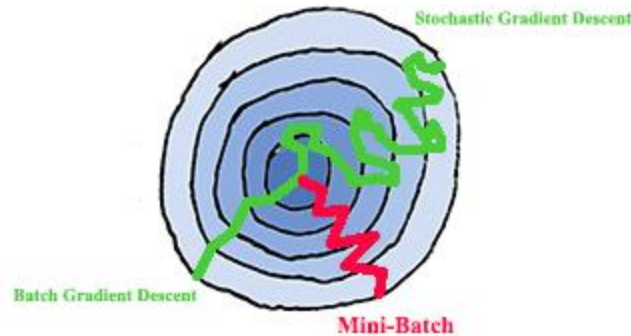
## Mini-batch Gradient descent





# Mini-batch and Epoch

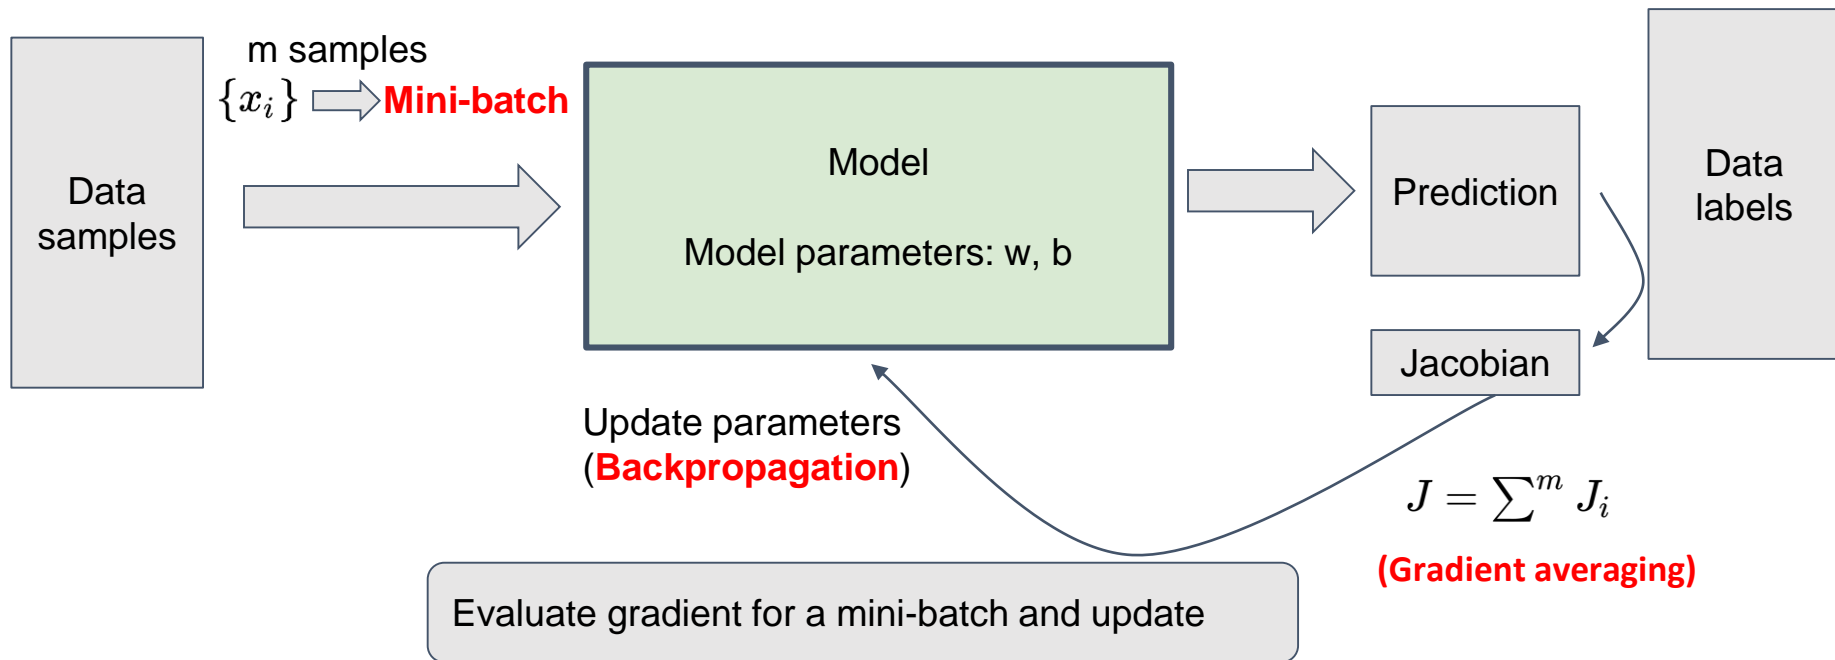
- **Mini-batch:** examples whose gradients are averaged before backpropagation
  - For more stable gradients



- **Epoch:** an ENTIRE dataset is passed forward and backward once
  - A complete training goes through several epoch

# Training in practice -- Overview

- For multiple epochs,....



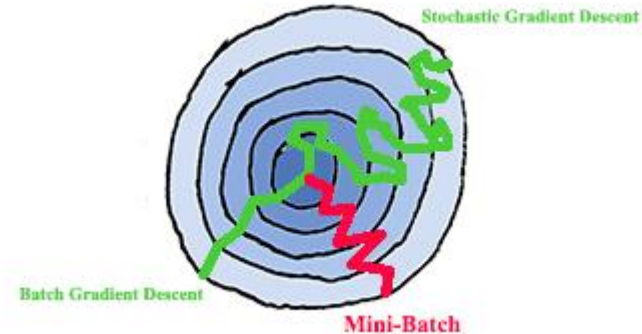


# Problems in Training?



# How good is SGD really?

- In the region of gentle slopes, gradients are very small.
- Near the minima loss are slightly ellipsoidal, where the first order Taylor approximations are almost orthogonal to the true direction of convergence.

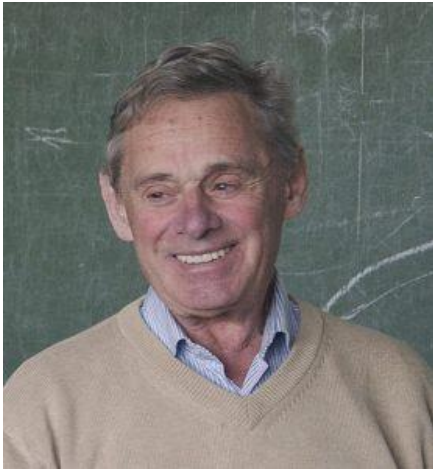


**Thought:** *Second order Taylor approximation may help.  
But using Newton (or Gauss-Newton) is too expensive.*



# Back in USSR

**Polyak momentum (1964)**



Boris Polyak

**Nesterov momentum (1983)**

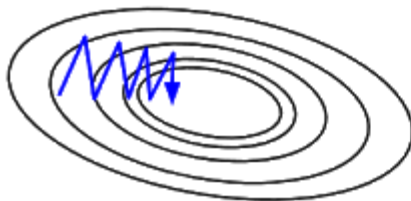


Yurii Nesterov

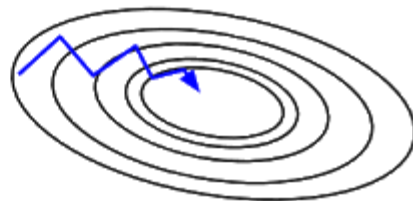
# Polyak Momentum

Momentum:  $\mathbf{m}_t = \underset{\text{Damping factor}}{\eta} \mathbf{m}_{t-1} + \underset{\text{Learning rate}}{\lambda} \nabla L(\mathbf{x}_t)$

Update:  $\mathbf{x}_{t+1} = \mathbf{x}_t - \mathbf{m}_t$



Stochastic Gradient Descent  
(SGD) without Momentum



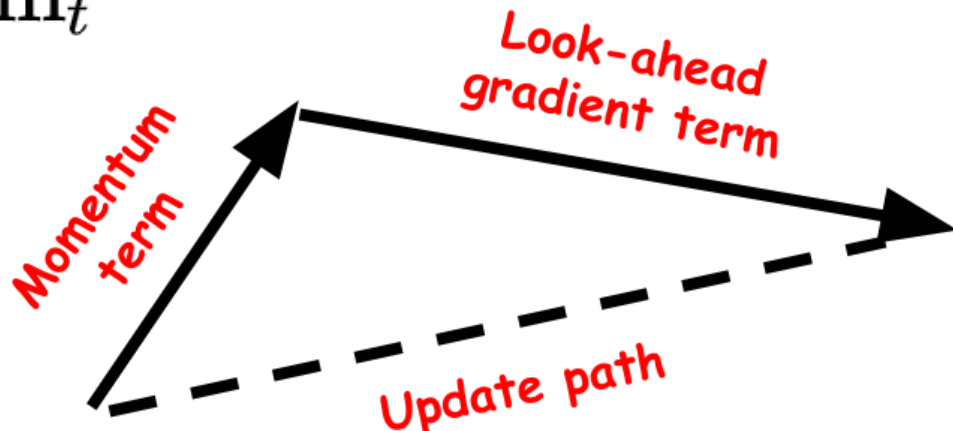
Stochastic Gradient Descent  
(SGD) with Momentum

# Nesterov Momentum

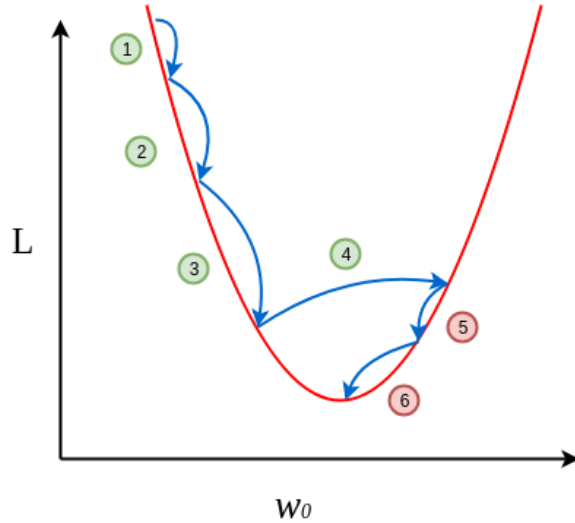
Look ahead:  $\mathbf{x}_{temp} = \mathbf{x}_t - \eta \mathbf{m}_{t-1}$

Momentum:  $\mathbf{m}_t = \eta \mathbf{m}_{t-1} + \lambda \nabla L(\mathbf{x}_{temp})$

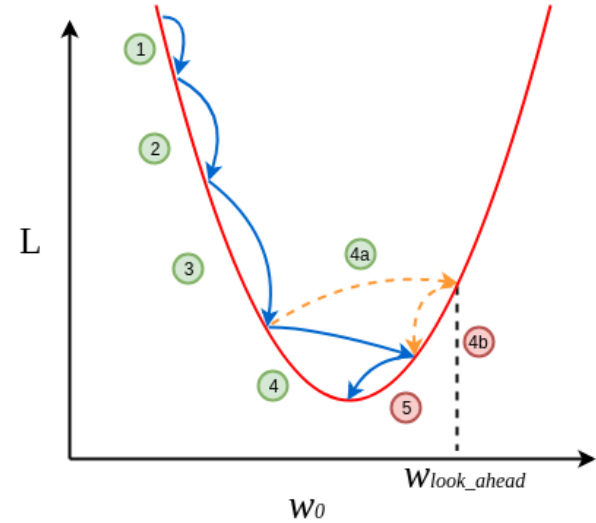
Update:  $\mathbf{x}_{t+1} = \mathbf{x}_t - \mathbf{m}_t$



# Polyak vs. Nesterov Momentum



(a) Momentum-Based Gradient Descent



(b) Nesterov Accelerated Gradient Descent

# Playing with the learning rate.....

- **AdaGrad:** larger updates for infrequent and smaller updates for frequent parameters

$$\mathbf{x}_t = \mathbf{x}_{t-1} - \frac{\lambda}{\text{RMS}[\nabla L(\mathbf{x}_t)]_t} \nabla L(\mathbf{x}_t)$$

Where  $\text{RMS}[\nabla L(\mathbf{x}_t)]_t$  is the sum of the squares of the past gradients.

```
torch.optim.Adagrad(params, lr=0.01, eps=1e-10)
```

- **Major issue:** rapid decay!



# Fixing AdaGrad.....

- **AdaDelta:** monotonically decrease using running average within some fixed window.

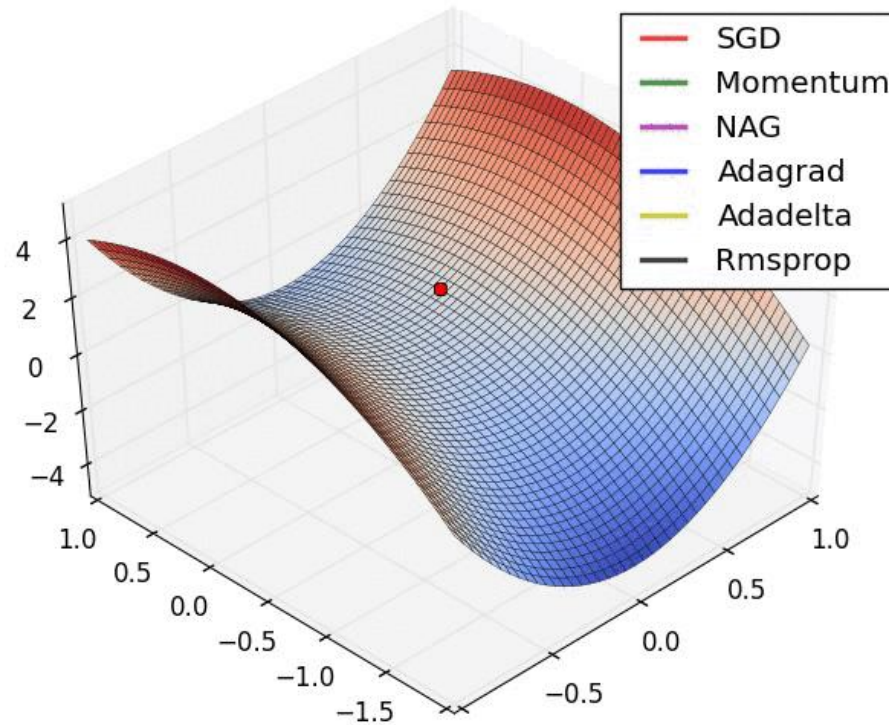
$$\mathbf{x}_t = \mathbf{x}_{t-1} - \frac{\text{RMS}[\Delta \mathbf{x}]_{t-1}}{\text{RMS}[\nabla L(\mathbf{x}_t)]_t} \nabla L(\mathbf{x}_t)$$

- **RMSProp:** decrease learning rate using a moving average of the squared gradients

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \frac{\lambda}{\sqrt{v_t + \epsilon}} \nabla L(\mathbf{x}_t) \quad \text{where,} \quad v_t = (1 - \alpha) \nabla L(\mathbf{x}_t)^2 + \alpha v_{t-1}$$

```
torch.optim.RMSprop(params, lr=0.01, alpha=0.99, eps=1e-08)
```

So, finally .....what are the possible choices?



# Adam: adaptive moment estimation

- Adam is a combination of RMSprop and SGD with momentum

## RMSprop

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \frac{\lambda}{\sqrt{v_t + \epsilon}} \nabla L(\mathbf{x}_t)$$

## SGD with momentum

$$\text{Momentum: } \mathbf{m}_t = \eta \mathbf{m}_{t-1} + \lambda \nabla L(\mathbf{x}_t)$$

$$\text{Update: } \mathbf{x}_{t+1} = \mathbf{x}_t - \mathbf{m}_t$$

$$\text{Update: } \mathbf{x}_{t+1} = \mathbf{x}_t - \frac{\lambda \hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$$

$$\text{where, } \hat{m}_t = \frac{m_t}{1 - \beta_1^t} \text{ and } \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

with,

$$m_t = (1 - \beta_1) \nabla L(\mathbf{x}_t) + \beta_1 m_{t-1}$$

$$v_t = (1 - \beta_2) \nabla L(\mathbf{x}_t)^2 + \beta_2 v_{t-1}$$

Moving averages of gradient and squared gradient

```
torch.optim.Adam(params, lr=0.001, betas=(0.9, 0.999), eps=1e-08)
```

# Adam: adaptive moment estimation

## Adam: A Method for Stochastic Optimization

<https://arxiv.org> › cs ▼

by DP Kingma - 2014 - Cited by 34535 - Related articles

Abstract: We introduce **Adam**, an algorithm for first-order gradient-based **optimization** of stochastic objective functions, based on adaptive estimates of ...

**Require:**  $\alpha$ : Stepsize

**Require:**  $\beta_1, \beta_2 \in [0, 1)$ : Exponential decay rates for the moment estimates

**Require:**  $f(\theta)$ : Stochastic objective function with parameters  $\theta$

**Require:**  $\theta_0$ : Initial parameter vector

$m_0 \leftarrow 0$  (Initialize 1<sup>st</sup> moment vector)

$v_0 \leftarrow 0$  (Initialize 2<sup>nd</sup> moment vector)

$t \leftarrow 0$  (Initialize timestep)

**while**  $\theta_t$  not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$  (Get gradients w.r.t. stochastic objective at timestep  $t$ )

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$  (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$  (Update biased second raw moment estimate)

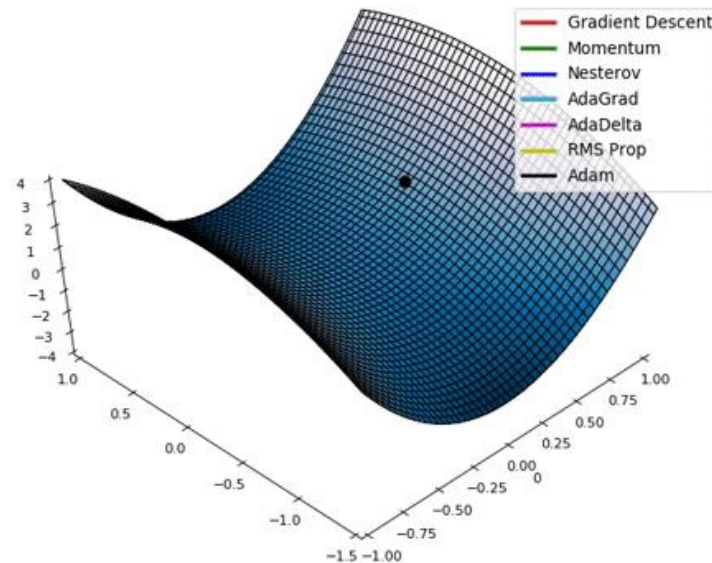
$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$  (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$  (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$  (Update parameters)

**end while**

**return**  $\theta_t$  (Resulting parameters)

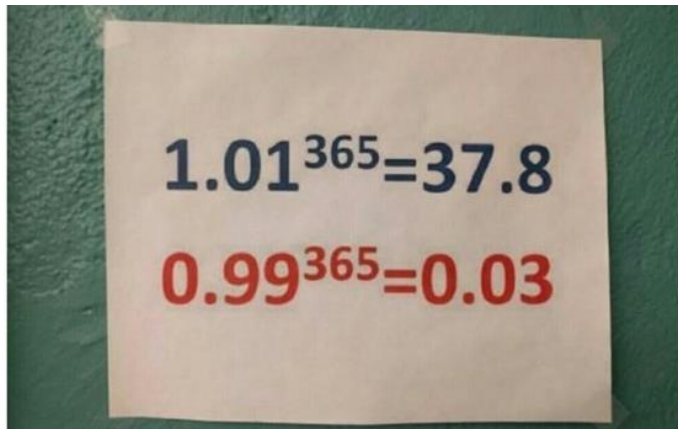


It's finally done! You understood everything.

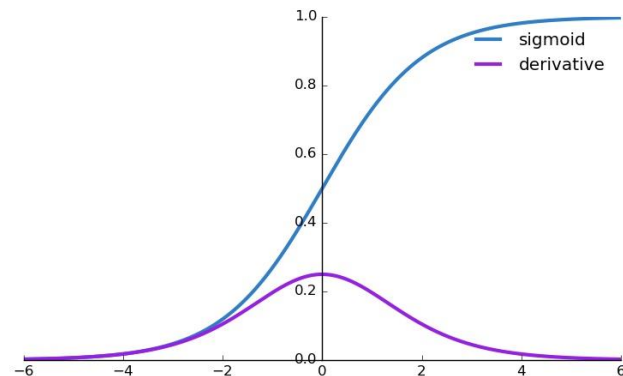




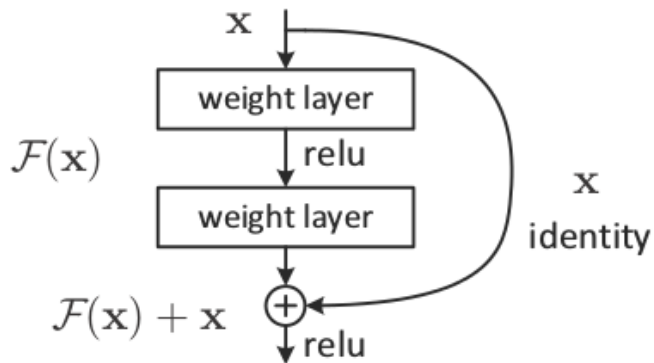
# Vanishing/Exploding gradient?



$1.01^{365}=37.8$   
 $0.99^{365}=0.03$

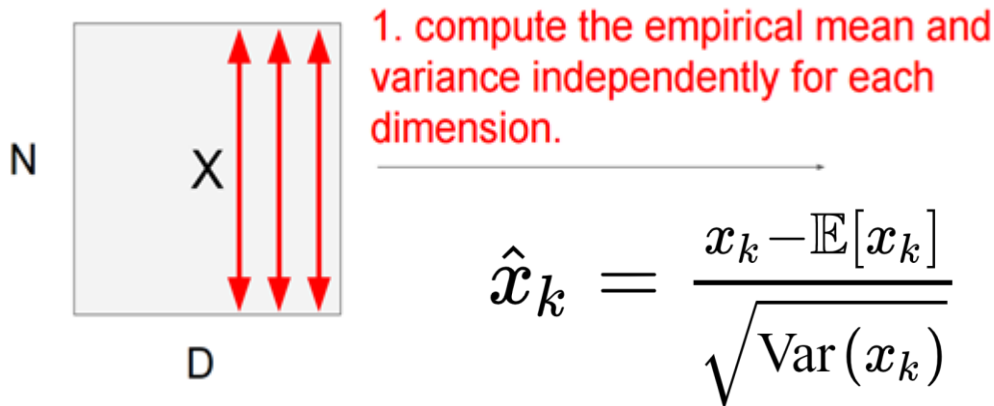


- Treatments
  - Weight initialization
  - Skip connections



# Need better gradients...??

## □ Batch normalization



```
nn.BatchNorm1d(input_size)
```



# Deep Learning Fundamentals -- Checklist\*

- Settings...

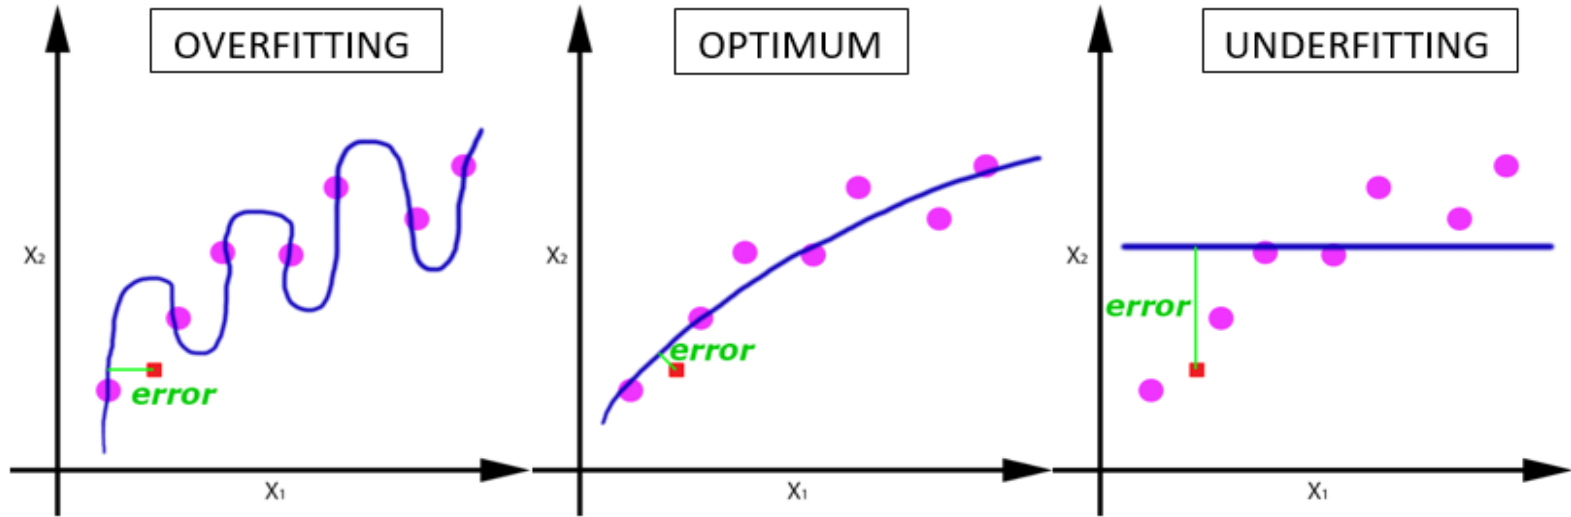
- Optimizers ✓
- Learning rate ✓
- Batch size ✓
- ~~Hyperparameters~~
- Weight decay
- Dropouts

- Design...

- Batch Normalization ✓
- ~~Layers/Channels~~
- Activations
- Data Augmentation
- Loss functions
- Softmax

\*non exhaustive!

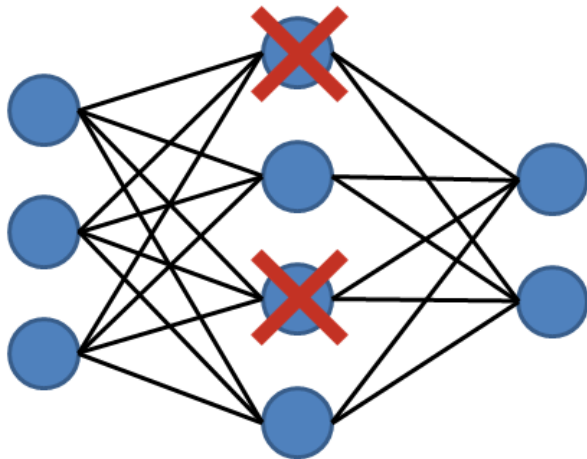
# What is Overfitting?



- Treatments:
  - Dropouts
  - Weight decay
  - Data augmentation

# Overfitting?

## ❏ Dropouts



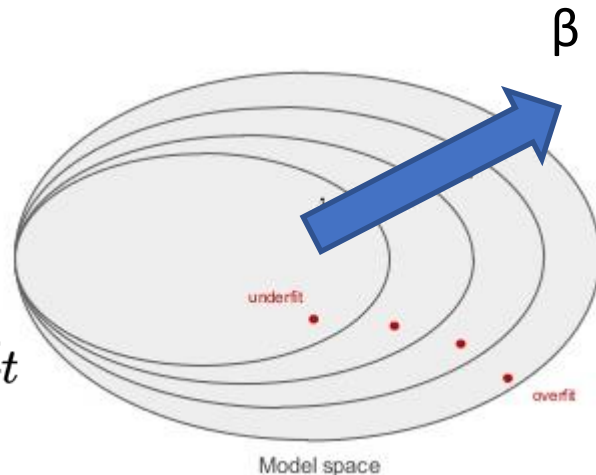
```
nn.Dropout (p=0.4)
```

# Overfitting??

- Weight decay:
  - Add penalty for large weights

$$L(\mathbf{x}_t) = L_{data}(\mathbf{x}_t) + \frac{\beta}{2} ||\mathbf{x}_t||^2$$

$$\text{Update: } \mathbf{x}_{t+1} = \mathbf{x}_t - \lambda \nabla L_{data}(\mathbf{x}_t) - \beta \mathbf{x}_t$$



```
torch.optim.Adam(params, lr=0.001, betas=(0.9, 0.999), eps=1e-08, weight_decay=0.01)
```

# Overfitting??

- ❑ Data Augmentation: make the problem more difficult



Data augmentation

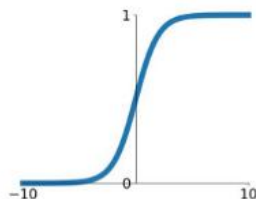


```
transforms.RandomHorizontalFlip(); transforms.CenterCrop(); ts.transforms.Rotate(); .....
```

# Activation functions: ReLu, Sigmoid, and more...

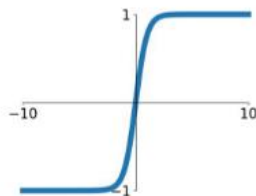
## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



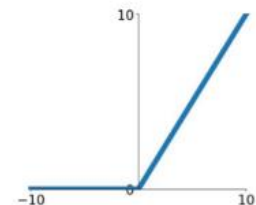
## tanh

$$\tanh(x)$$



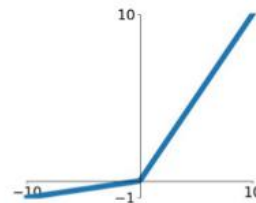
## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$

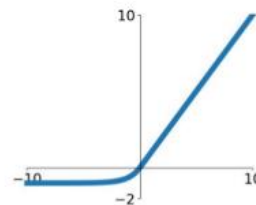


## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



```
nn.ReLU(); nn.Sigmoid(); nn.Tanh() .....
```

# What loss to minimize?

- L1-norm, L2-norm,
- Robust losses.....

$$\text{Least-square} := \|\hat{y} - y\|^2$$

Prediction

ground-truth

Name	$\rho(x)$	$\psi(x)$	$\omega(x)$
Least-squares	$x^2/2$	$x$	1
$L_1$ -norm	$ x $	$\text{sgn}(x)$	$1/ x $
$L_p$ -norm	$ x ^p/p$	$\text{sgn}(x) x ^{p-1}$	$ x ^{p-2}$
Fair	$\xi^2(\frac{ x }{\xi} - \log(1 + \frac{ x }{\xi}))$	$\frac{x}{1+ x /\xi}$	$\frac{1}{1+ x /\xi}$
Cauchy	$\frac{\xi^2}{2} \log(1 + x^2/\xi^2)$	$\frac{x}{(1+x^2/\xi^2)}$	$\frac{1}{(1+x^2/\xi^2)}$
Huber	$\begin{cases} x^2/2 &  x  \leq \xi \\ \xi( x  - \xi/2) &  x  > \xi \end{cases}$	$\begin{cases} x &  x  \leq \xi \\ \xi \text{sgn}(x) &  x  > \xi \end{cases}$	$\begin{cases} 1 &  x  \leq \xi \\ \xi/ x  &  x  > \xi \end{cases}$
Tukey	$\begin{cases} \frac{x^6}{6} - \frac{\xi^2 x^4}{2} + \frac{\xi^4 x^2}{2} &  x  \leq \xi \\ \frac{\xi^6}{6} &  x  > \xi \end{cases}$	$\begin{cases} x(\xi^2 - x^2)^2 &  x  \leq \xi \\ 0 &  x  > \xi \end{cases}$	$\begin{cases} (\xi^2 - x^2)^2 &  x  \leq \xi \\ 0 &  x  > \xi \end{cases}$

- Cross Entropy  $\rightarrow$  logistic outputs (e.g. classification)

- In binary classification, the cross-entropy is:

$$-(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$$

- If the number of classes  $M > 2$  (i.e. multiclass classification), it turns out to be:

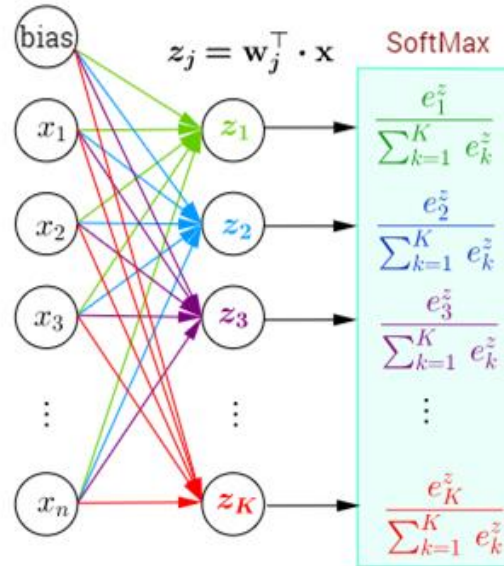
```
loss = nn.CrossEntropyLoss()
```

$$-\sum_{c=1}^M y_{o,c} \log(p_{o,c})$$

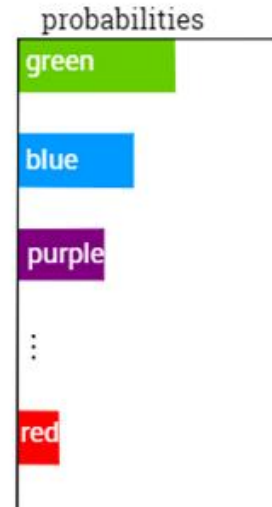
# Multi-class classification

## □ Softmax

$$\mathbf{z} = \begin{bmatrix} z_1 \\ z_2 \\ z_3 \\ \vdots \\ z_K \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^\top \\ \mathbf{w}_2^\top \\ \mathbf{w}_3^\top \\ \vdots \\ \mathbf{w}_K^\top \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \\ x_n \end{bmatrix}$$



Multiple outputs of probabilities

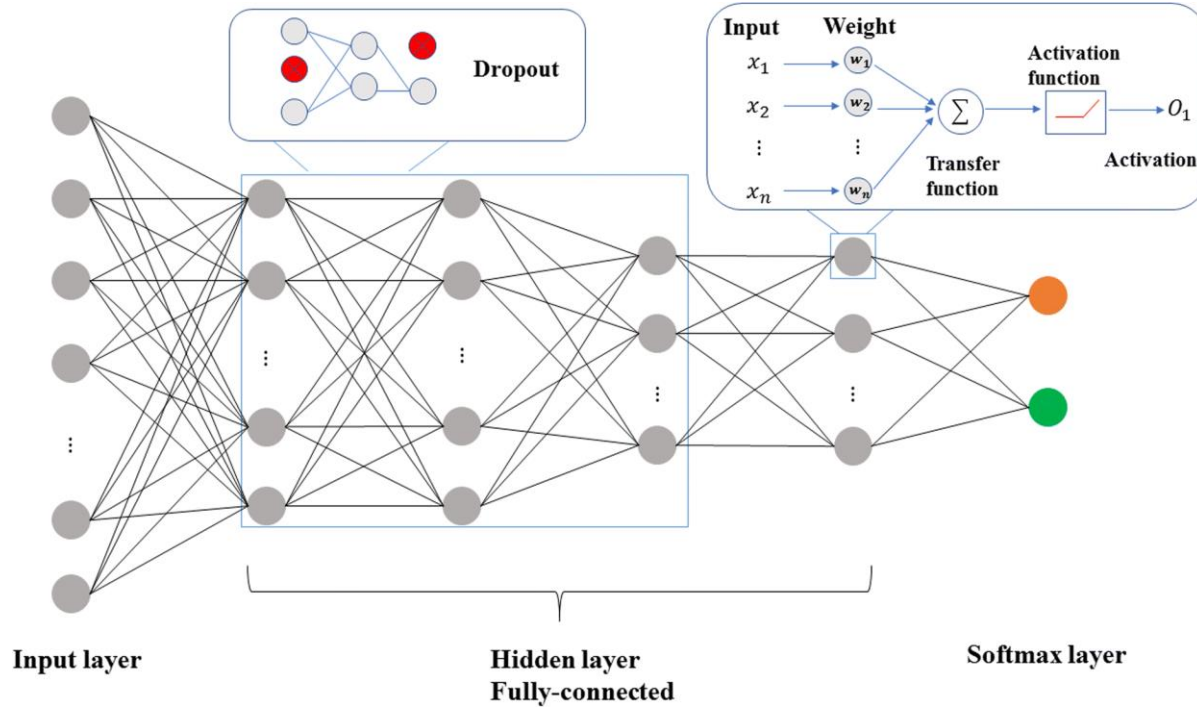


```
nn.functional.softmax(input)
```





# A General Overview.....



# Lab session...

- Install *pytorch*
- Plot the loss function and accuracy evolution with change in
  - Learning rate
  - Batch size
  - Optimizers

Using .... github code: [yunjey/pytorch-tutorial](https://github.com/yunjey/pytorch-tutorial)

Link: [https://github.com/yunjey/pytorch-tutorial/blob/master/tutorials/01-basics/logistic\\_regression/main.py](https://github.com/yunjey/pytorch-tutorial/blob/master/tutorials/01-basics/logistic_regression/main.py)

- Feel free to try other “[pytorch-tutorial](#)” code in the same repository!